



PayPal Merchant ecosystem using Spark, Hive, Druid, HBase & Elasticsearch

Deepika Khera & Kasiviswanathan Natarajan

June 19 2018

Who we are?



Deepika Khera

- Big Data Technologist for over a decade
- Focused on building scalable platforms with Hadoop ecosystem – Map Reduce, HBase, Spark, Elasticsearch, Druid
- Senior Engineering Manager - Merchant Analytics at PayPal
- Contributed to Druid for the Spark Streaming integration



Kasi Natarajan

- 15+ years of industry experience
- Spark Engineer @PayPal Merchant Analytics
- Building solutions using Apache Spark, Scala, Hive, HBase, Druid and Spark ML.
- Passionate about providing Analytics at scale from Big Data platforms

Agenda

- ❖ PayPal Data & Scale
- ❖ Merchant Use Case Review
- ❖ Data Pipeline
- ❖ Learnings - Spark & HBase
- ❖ Tools & Utilities
 - ❖ Behavioral Driven Development
 - ❖ Data Quality Tool using Spark
 - ❖ BI with Druid & Tableau

PayPal Data & Scale

PayPal is more than a button



CBT



Mobile



In-Store



Online



Loyalty



Credit



APV Lift



Offers



Faster
Conversion



Reduction
in Cart
Abandonment



Customer
Acquisition

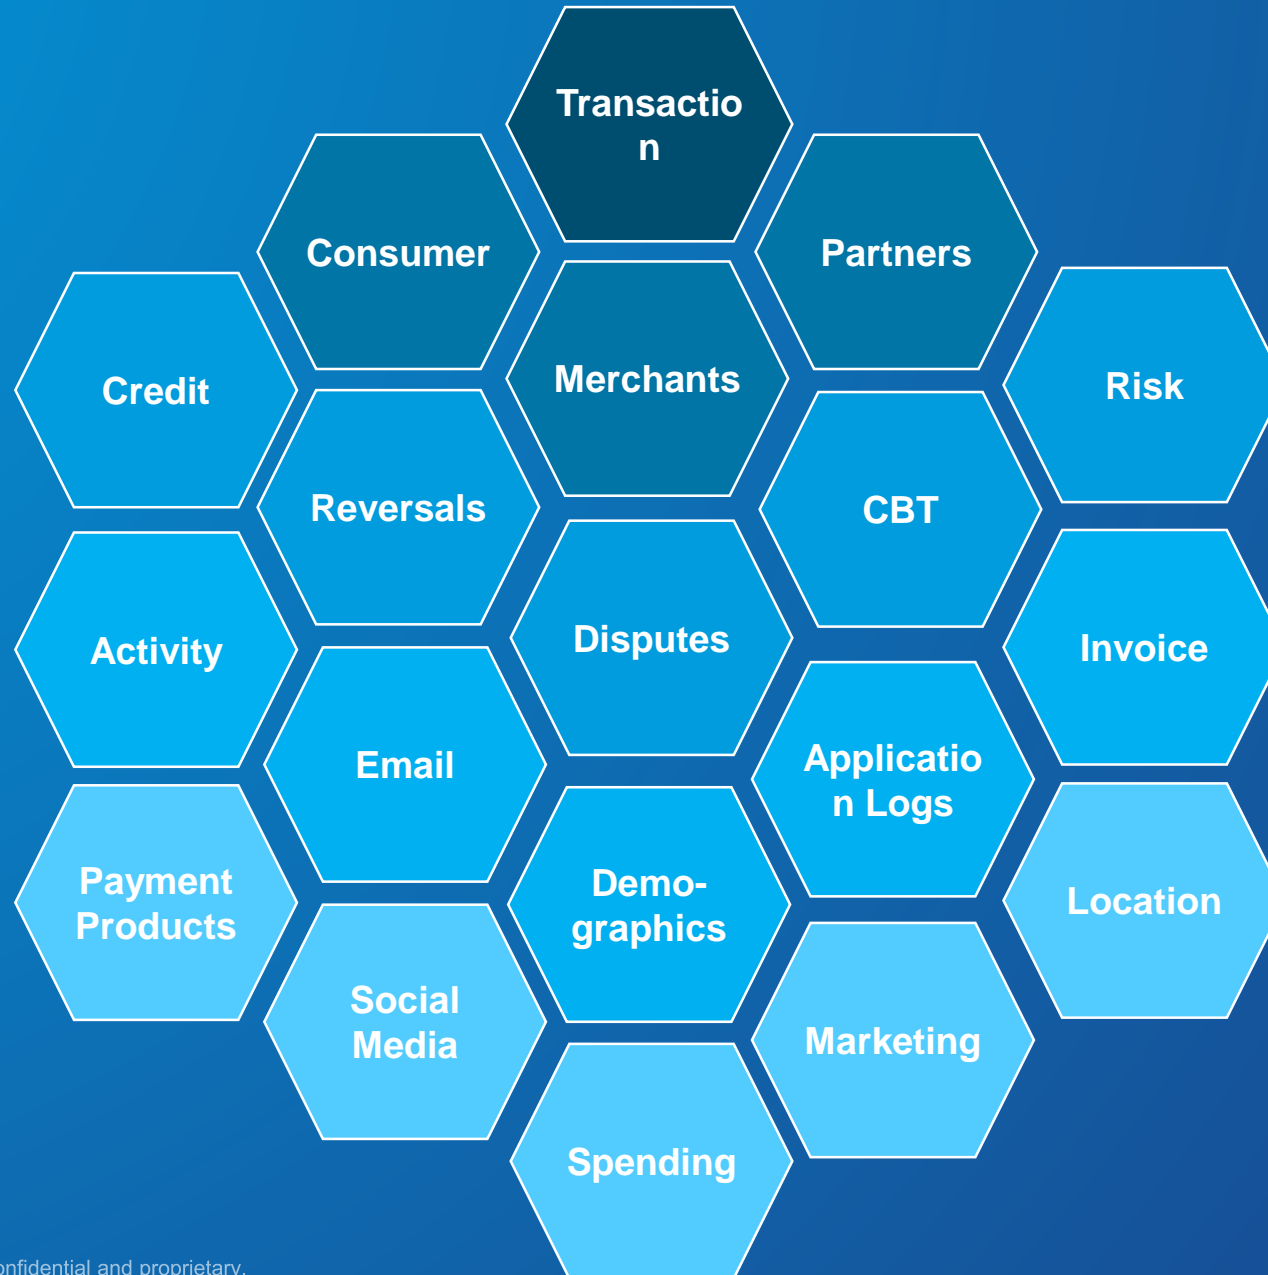


Invoicing

Check out with **PayPal**



PayPal Datasets



The power of our platform

PayPal operates one of
the largest

**PRIVATE
CLOUDS**
in the world*

7.6
BILLION
payments in
2017**

~60
payments/
second at peak*


19M
merchants

200
+
markets

237M
active customer
accounts**

42
petabytes
of data*

Dedicated to with a
customer focused,
strong performance,
highly scalable,
continuously available
PLATFORM.



PayPal operates one of the largest Hadoop
deployments in the world.
A 1600 Node Hadoop Cluster
with 230TB of Memory, 78PB of Storage
Running 50,000 Jobs Per day

PayPal has one of the top five Kafka
deployments in the world, handling over
200 billion messages per day

Merchant Use Case Review

Use Case Overview



INSIGHTS



- Revenue & transaction trends
- Cross-Border Insights
- Customer Shopping Segments



MARKETING SOLUTIONS

- Help Merchants engage their customers by personalized shopping experience
- Offers & Campaigns
- Shoppers Insights



PAYPAL ANALYTICS

- Products performance
- Checkout Funnel
- Behavior analysis
- Measuring effectiveness

Merchant Data Platform

1. Fast Processing platform crunching multi-terabytes of data
2. Scalable, Highly available, Low latency Serving platform

Technologies



Processing



elastic



druid

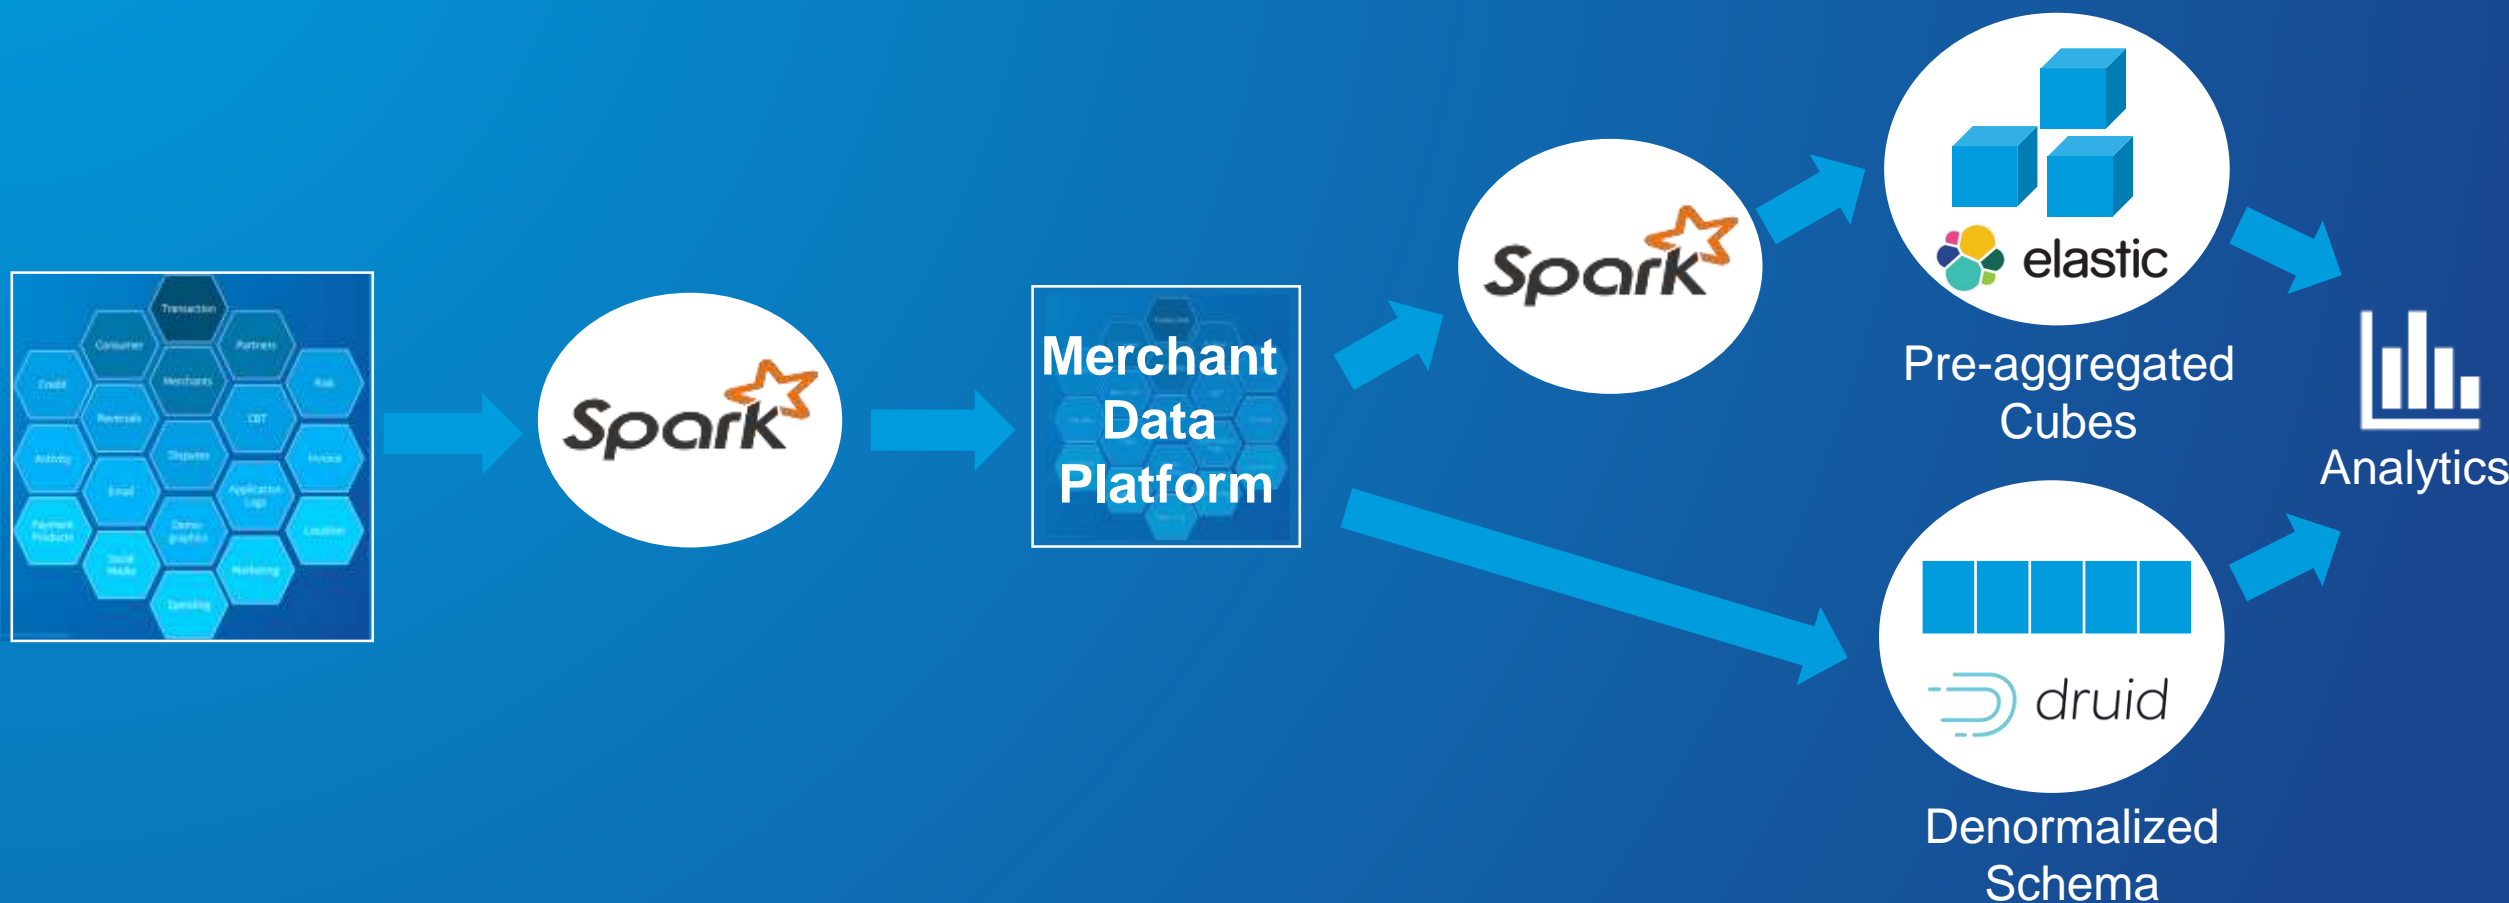
Serving



kafka

Movement

Merchant Analytics



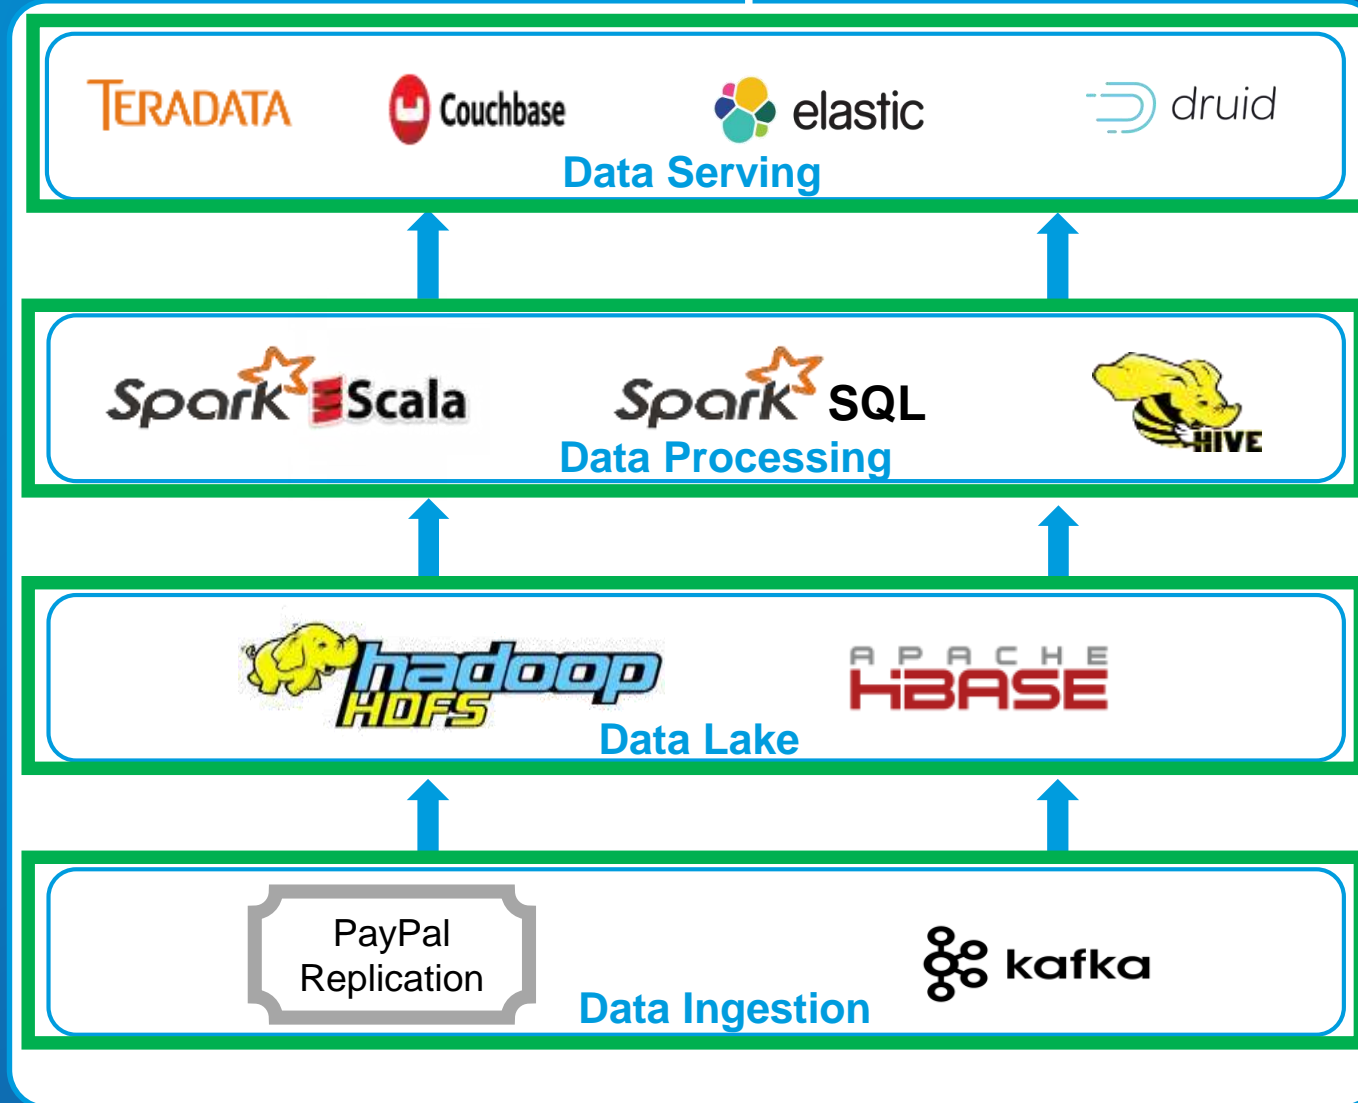
Data Pipeline

Data Pipeline Architecture


Data Sources



Data Pipeline



Visualization



Learnings – Spark & HBase

Spark Best Practices Checklist



Data Serialization

- ✓ Use Kryo Serializer with SparkConf, which is faster and compact
- ✓ Tune kryoserializer buffer to hold large objects



Garbage Collection

- ✓ Clean up cached/persisted collections when they are no longer needed
- ✓ Tuned concurrent abortable preclean time from 10sec to 30sec to push out stop the world GC



Memory Management

- ✓ Avoided using executors with too much memory



Parallelism

- ✓ Optimize number of cores & partitions*



Action-Transformation

- ✓ Minimize shuffles on join() by broadcasting the smaller collection
- ✓ Optimize wider transformations as much as possible*



Caching & Persisting

- ✓ Used MEMORY_AND_DISK storage level for caching large
- ✓ Repartition data before persisting to HDFS for better performance in downstream jobs

*Specific examples later

Learnings

Spark job failures with Fetch Exceptions



Observations

- Executor spends long time on shuffle reads. Then times out, terminates and results in job failure
- Resource constraints on executor nodes causing delay in executor node

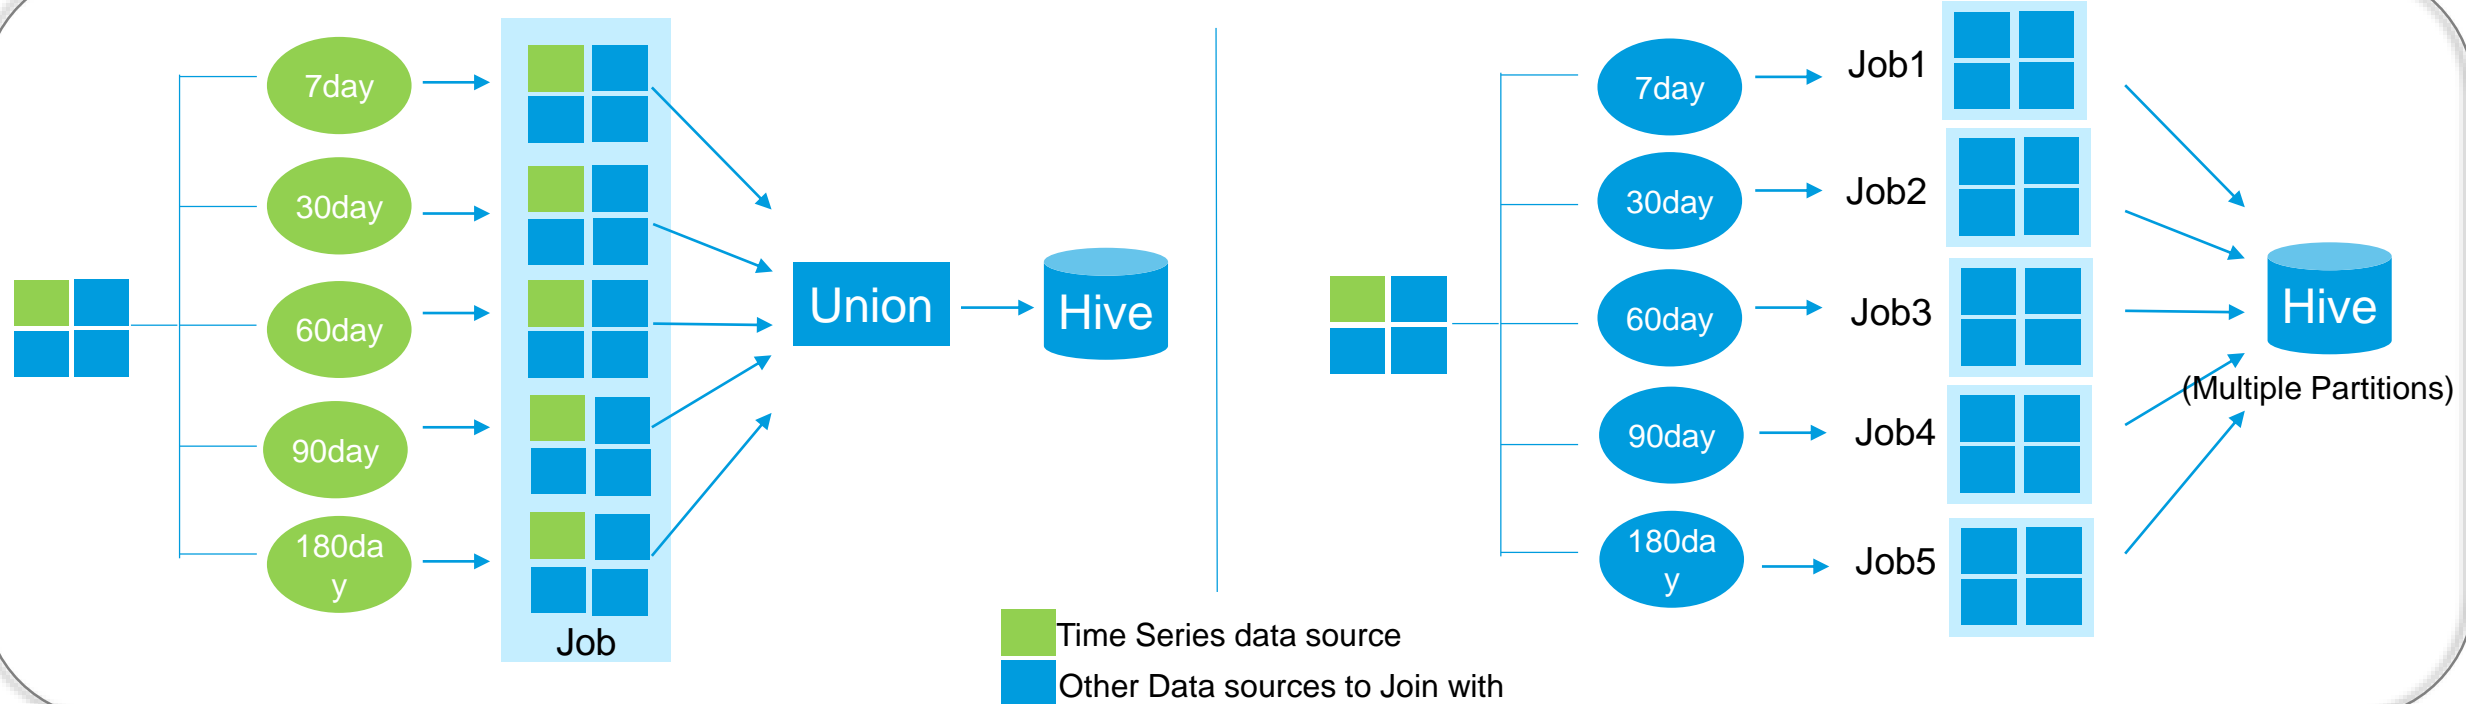
Resolution

To address memory constraints, tuned

1. config from 200 executor * 4 cores to 400 executor * 2 cores
2. executor memory allocation (reduced)

Learnings

Parallelism for long running jobs



Observations

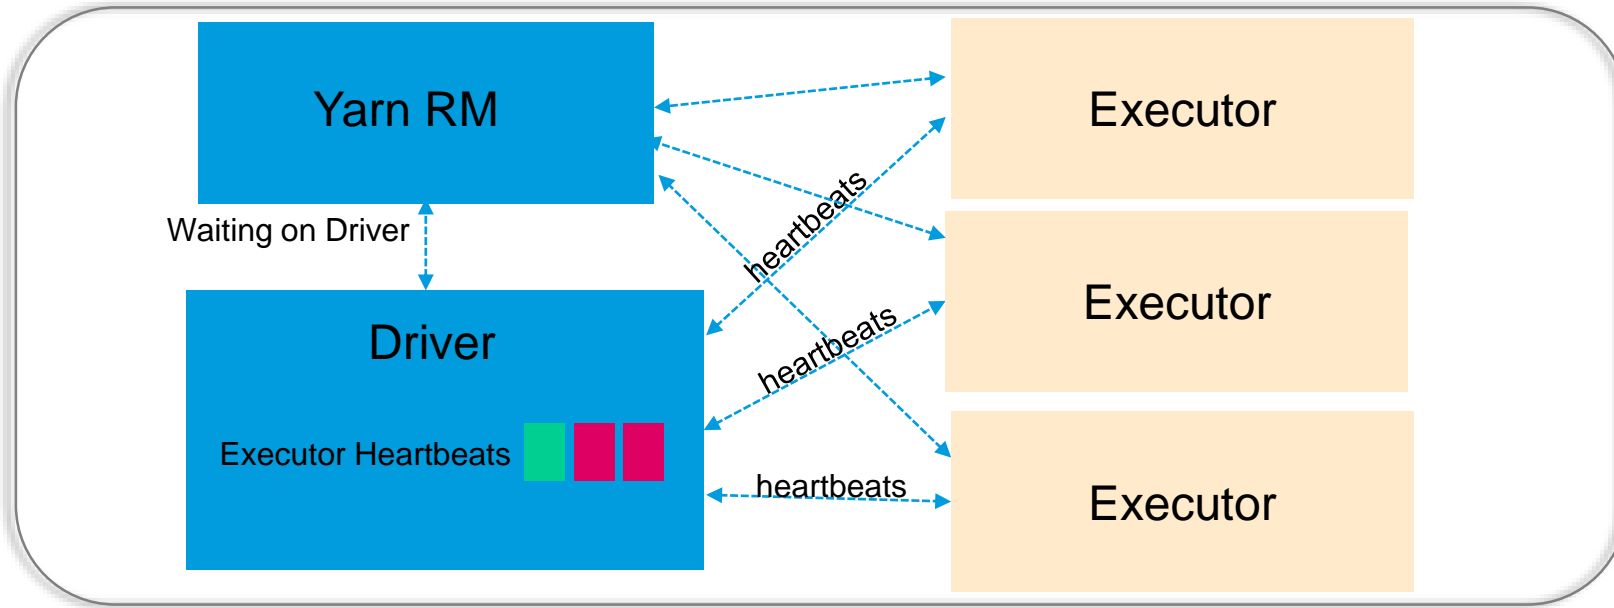
- Series of left joins on large datasets cause shuffle exceptions

Resolution

1. Split into Small jobs and run them in parallel
2. Faster reprocessing and fail fast jobs

Learnings

Tuning between Spark driver and executors



Observations

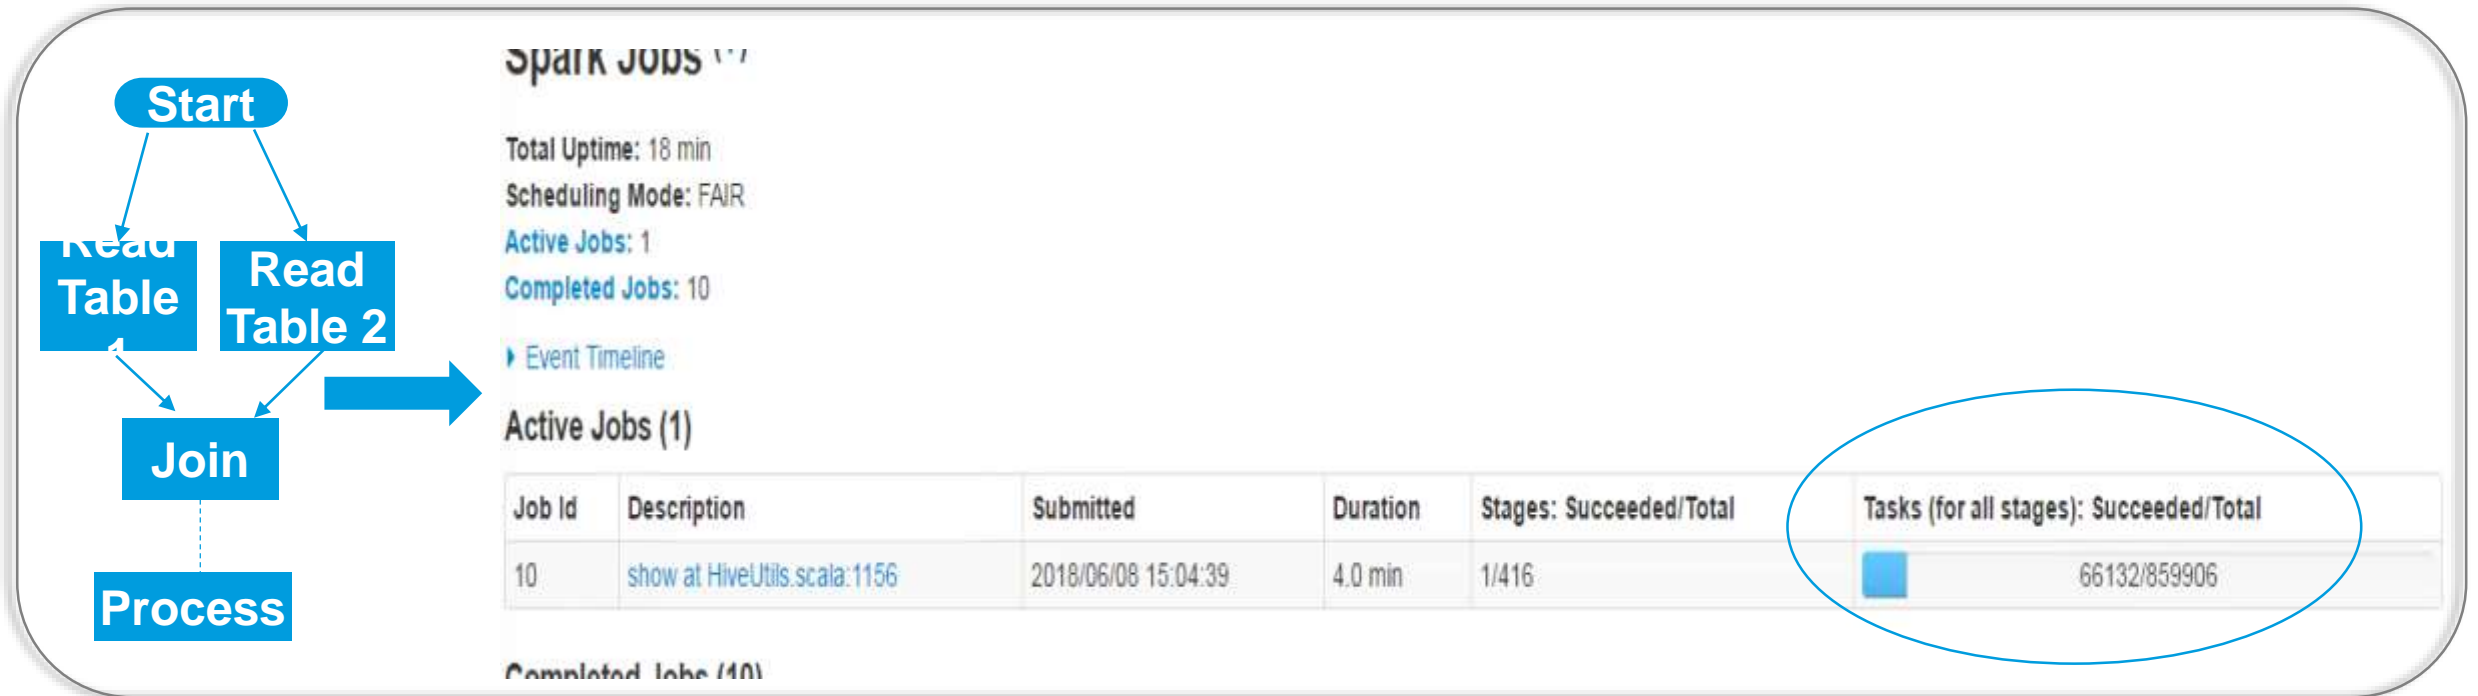
- Spark Driver was left with too many heartbeat requests to process even after the job was complete
- Yarn kills the Spark job after waiting on the Driver to complete processing the Heartbeats

Resolution

- The setting "spark.executor.heartbeatInterval" was set too low. Increasing it to 50s fixed the issue
- Allocate more memory to Driver to handle overheads other than typical Driver processes

Learnings

Optimize joins for efficient use of cluster resources (Memory, CPU etc..,)



Observation

With the default shuffle partitions of 200, the Join Stage was running with too many tasks causing performance overhead


Resolution

Reduce the `spark.sql.shuffle.partitions` settings to a lower threshold


Learnings

Optimize wide transformations

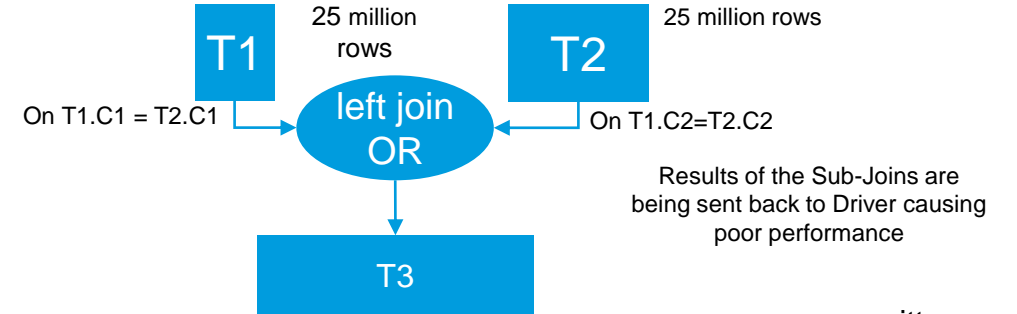
Left Outer Join




rewritten as



Left Outer Join with OR Operators



rewritten as



Resolution

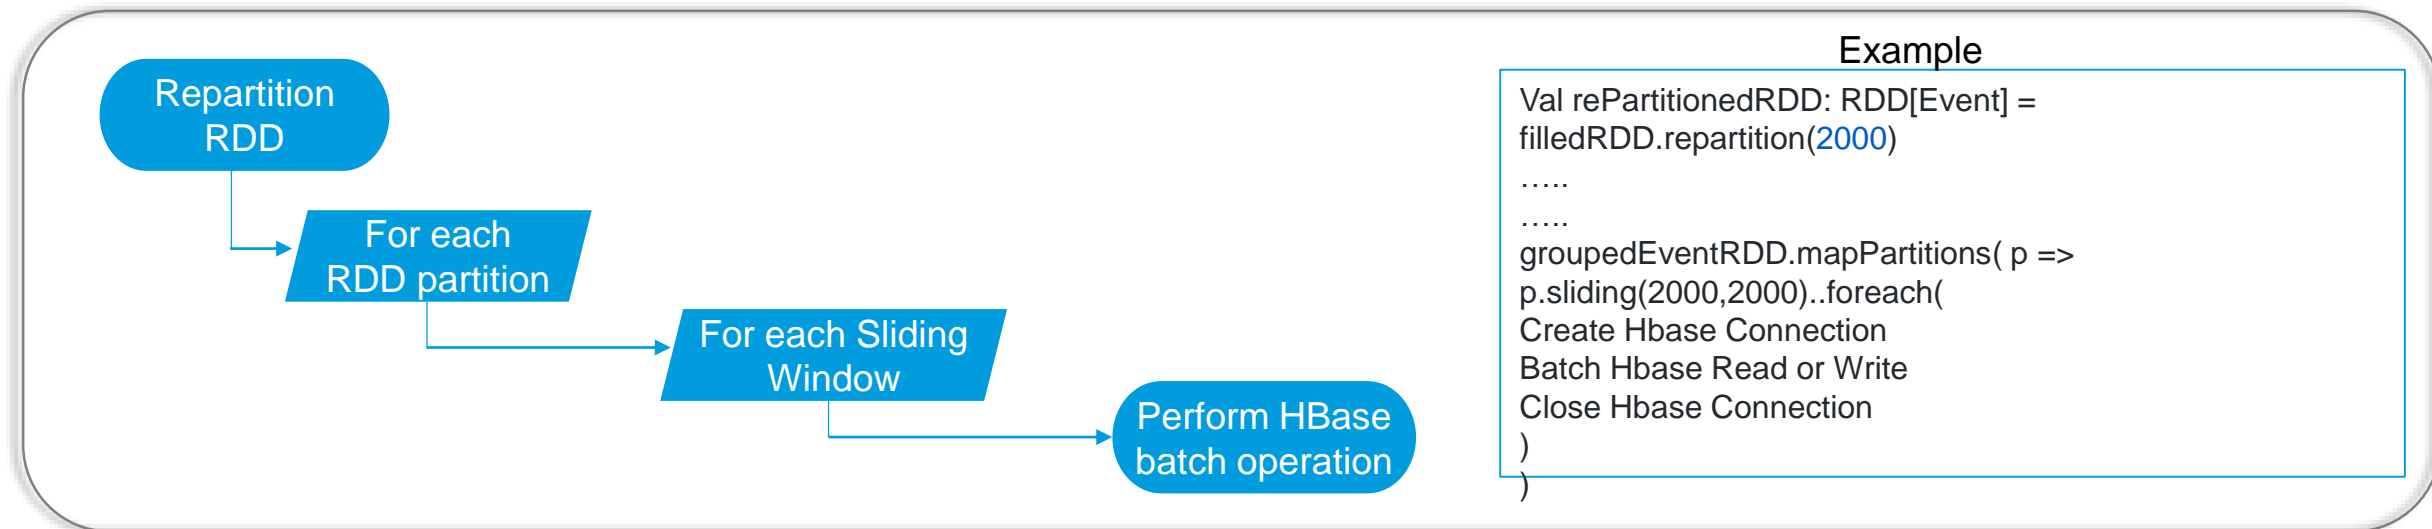
- Convert expensive left joins to combination of light weight join and except/union etc.,

Learnings

Optimize throughput for HBase Spark Connection

Observations

- Batch puts and gets slow due to HBase overloaded connections
- Since our HBase row was wide, HBase operations for partitions containing larger groups were slow




Resolution

- Implemented sliding window for HBase Operations to reduce HBase connection overload

Tools & Utilities

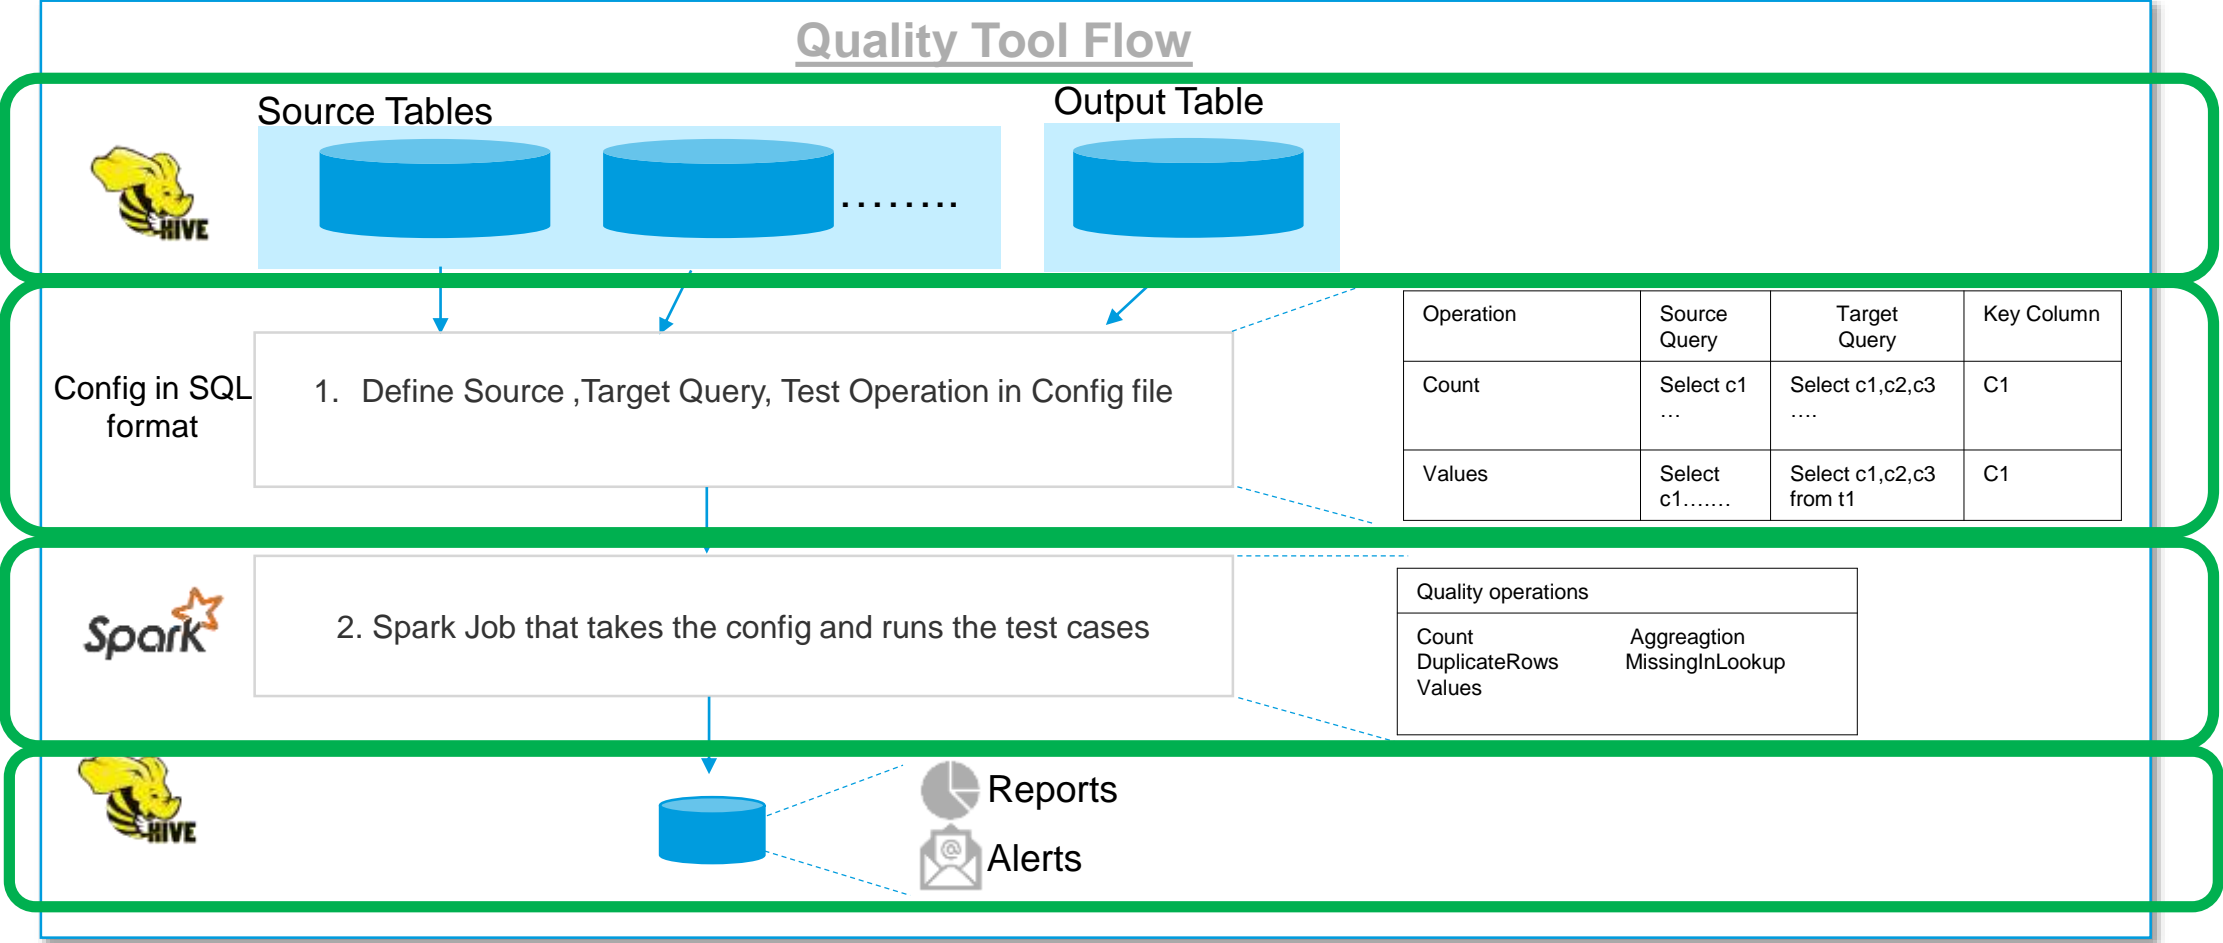
Behavioral Driven Development

- While Unit tests are more about the implementation, BDD emphasizes more on the behavior of the code
- Writing “Specifications” in pseudo-English.
- Enables testing at external touch-points of your application

Feature : Identify the activity related to an event	<u>pseudo code</u>
Scenario: Should perform an iteration on events and join to activity table and identify the activity name	<pre>import cucumber.api.scala.{EN, ScalaDsl} import cucumber.api.DataTable import org.scalatest.Matchers</pre>
Given I have a set of events cookie_id:String page_id:String last_actvty:String 263FHFBCBCCBV login_provide review_next_page HFFDJFLUFBFNJL home_page provide_credent	<pre>Given("""^I have a set of events\$""") { (data:DataTable) => eventdataDF = dataTableToDataFrame(data) }</pre>
And I have a Activity table last_activity_id:String activity_id:String activity_name:String review_next_page 1494300886856 Reviewing Next Page provide_credent 232323232323 Provide Credentials	 <pre>Given("""^I have a Activity table\$""") { (data:DataTable) => activityDataDF = dataTableToDataFrame(data) }</pre>
When I implement Event Activity joins	<pre>When("""^I implement Event Activity joins\$"""){ () => eventActivityDF = Activity.findAct(eventdataDF, activityDataDF) }}</pre>
Then the final result is cookie_id:String activity_id:String activity_name:String last_activity_id:String activity_id:String activity_name:String 263FHFBCBCCBV 1494300886856 Reviewing Next Page HFFDJFLUFBFNJL 232323232323 Provide Credentials	<pre>Then("""^the final result is \$"""){ (expectedData:DataTable) => val expectedDf = dataTableToDataFrame(expectedData) val resultDF = eventActivityDF resultDF.except(expectedDF).count</pre>

Data Quality Tool

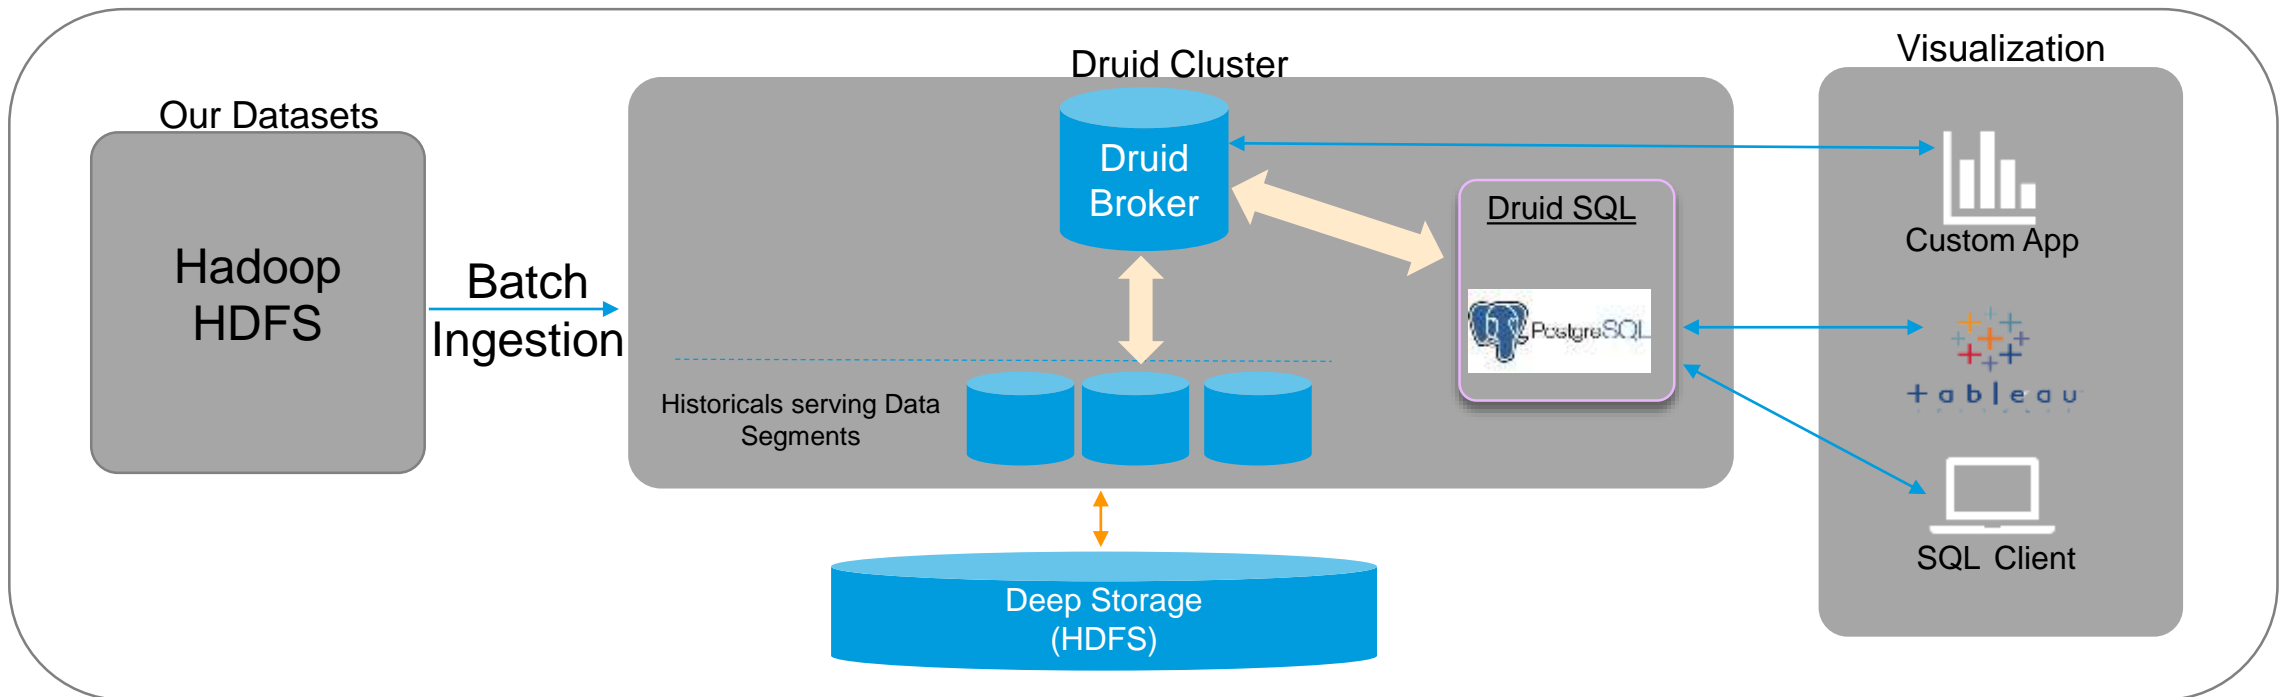
- Config driven automated tool written in Spark for Quality Control
- Used extensively during functional testing of the application and once live, used as quality check for our data pipeline
- Feature to compare tables (schema agnostic and at Scale) for data validation and helping engineers troubleshoot effectively



Druid Integration with BI

- Druid is an open-source time series data store designed for sub-second queries on real-time and historical data. It is primarily used for business intelligence queries on event data*
- Traditional Databases did not scale and perform with Tableau dashboards (for many use cases)
- Enable Tableau dashboards with Druid as the serving platform
- Live connection from tableau to druid avoids getting limited by storage at any layer.

Visualization at scale



Conclusion

- ❖ Spark Applications on Yarn (Hortonworks distribution).
- ❖ Spark jobs were easy to write and had excellent performance (though little hard to troubleshoot)
- ❖ Spark-HBase optimization improved performance
- ❖ Pre-aggregated datasets to Elasticsearch
- ❖ Denormalized datasets to Druid
- ❖ Pushed lowest-granularity denormalized datasets to Druid
- ❖ Behavior Driven Development a great add-on for Product-backed applications



QUESTIONS?